

CSCC09F

Programming on the Web



Introduction to the World Wide Web
(aka WWW, the Web)

Web Origins

- ❑ In 1989, Tim Berners-Lee working at CERN proposed an Intranet supporting document access within a heterogeneous environment
- ❑ TBL had an interest in hypertext, dating back to childhood, when he was exposed to a then-popular book, “Enquire Within”, that contained cross-reference “links” to many different topics – a sort of common-man’s how-to guide



ENQUIRE WITHIN

UPON

EVERYTHING.

“WHETHER YOU WANT TO KNOW A FLOWER IS DEAD; TO STUDY THE BEHAVIOUR OF BRIGITTE; TO WRITE A RECEIPT FOR BREAKFAST OR DINNER; TO READ A REVIEW FOR A BOOK; PARTS OF A SMALL ONE; TO CURSE A RENAISSANCE; TO MAKE A WILL; TO GET MARRIED; TO HAVE A RELATIVITY; WHATEVER YOU WANT TO DO, READ, OR TO KNOW, PROVIDE YOUR KNOWLEDGE HAS RELATION TO THE KNOWLEDGE OF HENRIETTA LEE, I SEND YOU WHAT YOU WANT TO ‘ENQUIRE WITHIN.’” — ZOLA.

ENQUIRERS ARE REFERRED TO THE INDEX AT THE END

EIGHTY-NINTH EDITION, REVISED

MAKING THE TOTAL SALES TO DATE


ONE MILLION ONE HUNDRED AND EIGHTY-EIGHT THOUSAND COPIES

LONDON:


HOULSTON AND SONS,

FATHERNOSTER SQUARE,

1894.



Modest Initial Goals

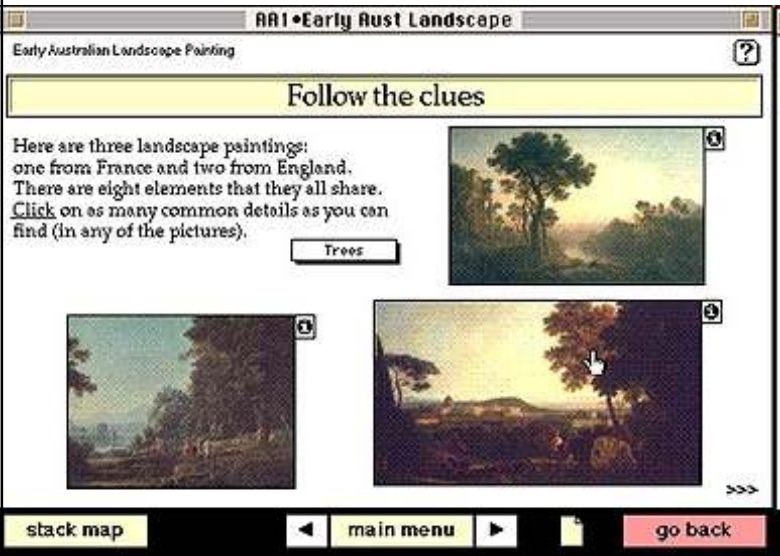


01 - World Wide Web

CSCC09

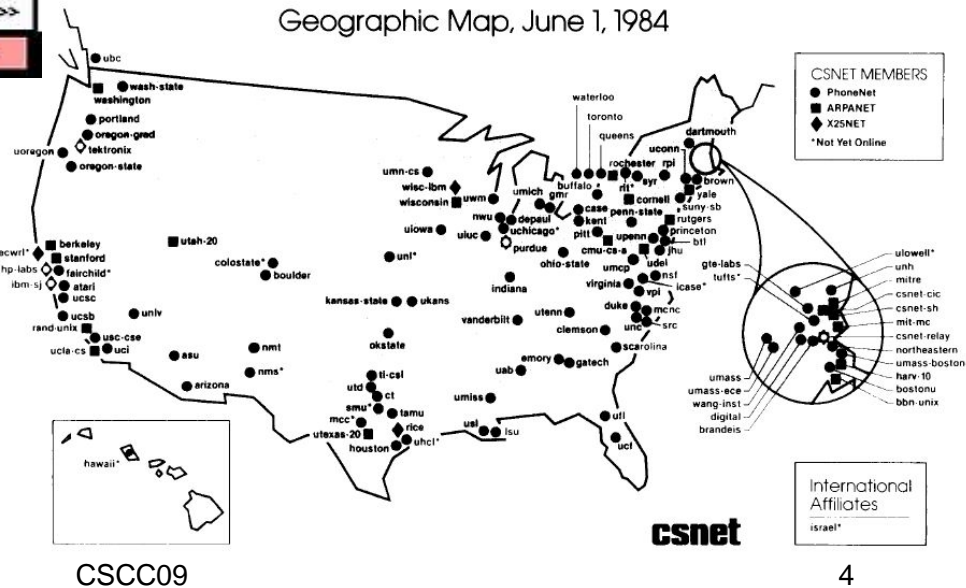
3

Web Origins: Groundwork



- Hypertext on computers dates back to '69, but especially work in '85 at Xerox (NoteCards) and '87 at Apple (HyperCard) which popularized the technique. Links were not networked.

- Meanwhile, the Internet was still a sleeper research & education network, unknown to the general population.




Web Origins: Groundwork

- ❑ TBL created a paradigm-shifting innovation, with considerable groundwork laid by earlier technologies:
 - SGML – a standardized “markup language” for describing document content, in use at CERN
 - hypertext – a way to link documents
 - the Internet – a standardized network framework with many existing protocols, e.g. email, ftp, telnet
 - Unix – with its regular, hierarchical file structure
 - and a large supporting cast, such as: client-server architecture, remote procedure calls (which TBL worked on at CERN), scripting languages, relational databases, etc.



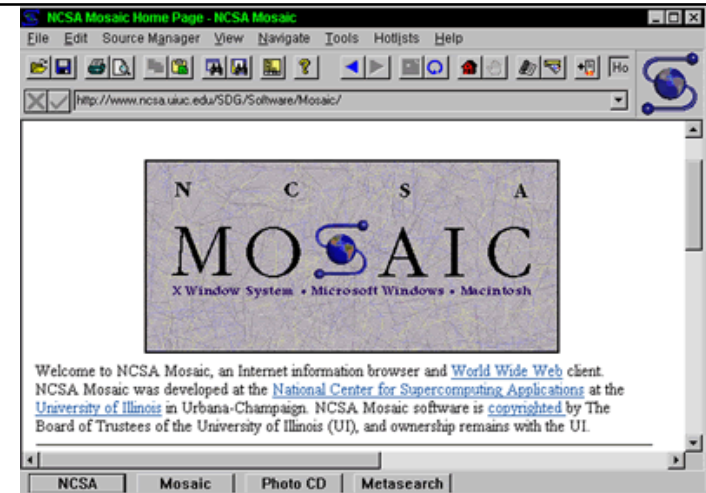
Web Origins: Tim's Synthesis



- ❑ So, what exactly did TBL invent?
 - defined and implemented a small markup language, based on an SGML application in use at CERN, which he named HTML
 - added network-traversable hyperlinks
 - operating over a new Internet application-protocol named HTTP
 - with a uniform naming structure (URLs) based on domain names from the Internet DNS and document locations specified as in Unix file paths
 - (he also implemented his ideas as a working system on a NeXT computer running a version of Unix)
- ❑ In hindsight it may all seem obvious, but it took a stroke of insight to pull all the existing pieces together into a coherent and compelling package

Web Origins: Out of the Lab

- ❑ quantum leap in Web use began with NCSA Mosaic browser:
 - image (img) tag; a multi-platform, easy-to-install and free GUI implementation -- the Web for non-geeks!
 - implemented first for X11 on UNIX(!), ports to Mac, Windows followed shortly
- ❑ evolved into a vehicle for solving a large number of important IT problems
- ❑ hey, we could make money on this ... Netscape Communications kicks off the Internet gold rush aka the “.com” boom!



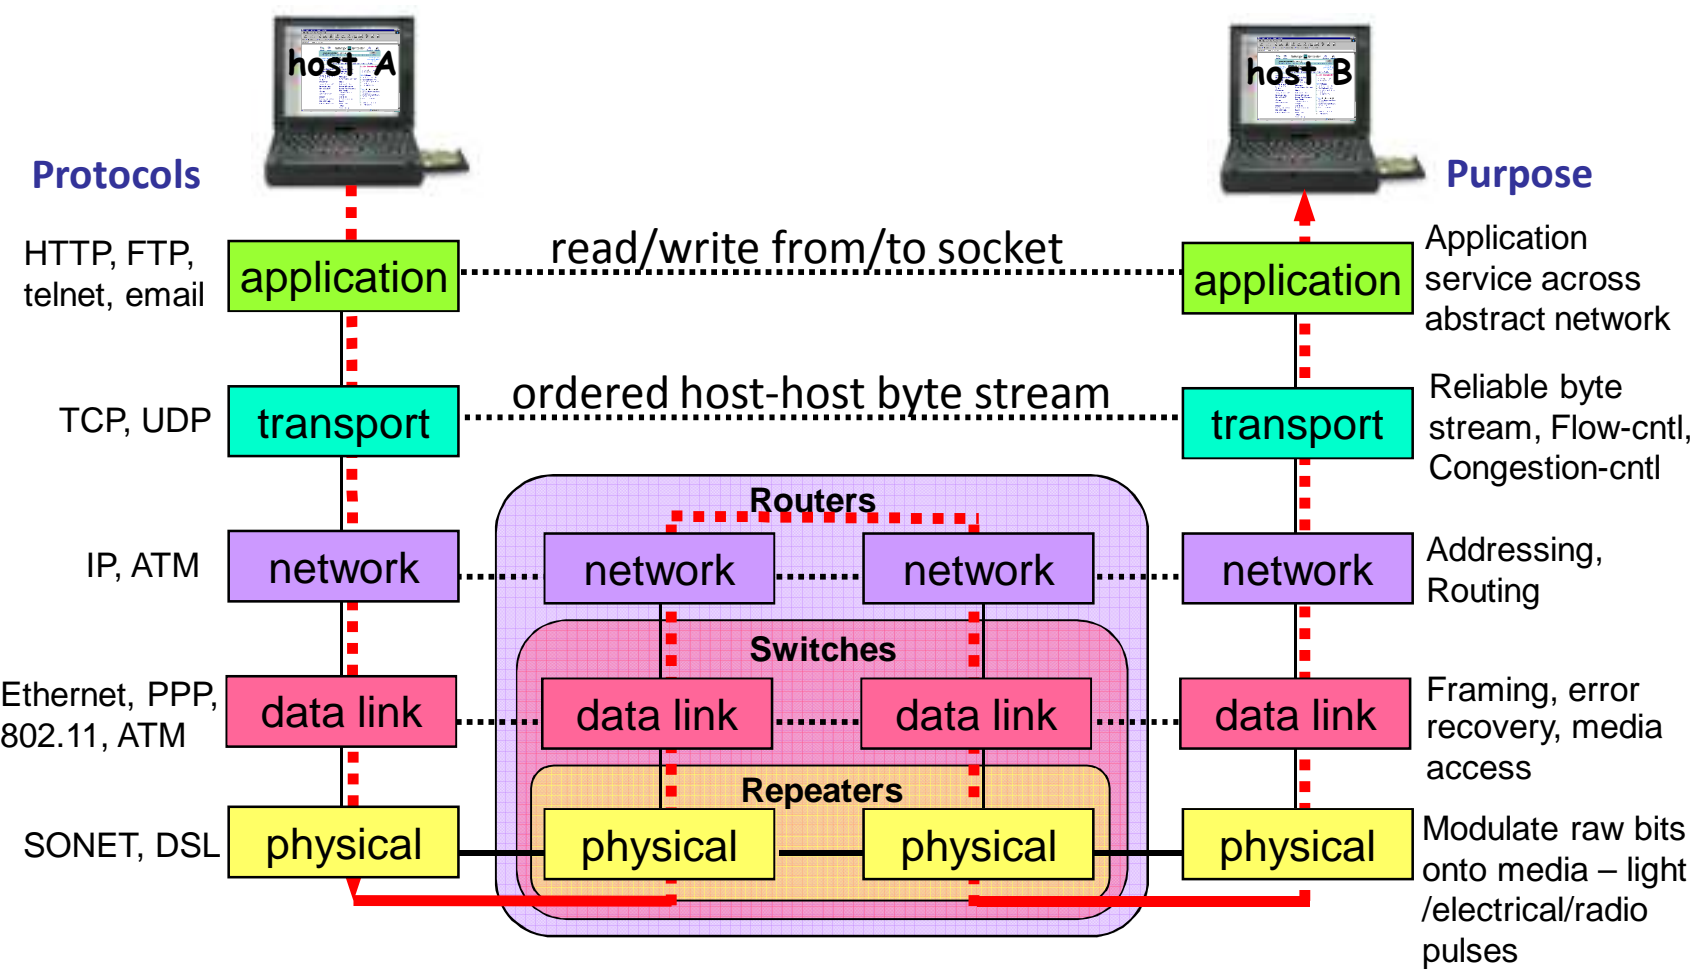
Web Components

- ❑ A collection of markup languages for describing the structure & meaning of document-content together with style descriptors to control document presentation
- ❑ A collection of protocols for moving content around the Net
 - especially HTTP (Hyper-Text Transfer Protocol), but also XHR, SSL/HTTPS, SOAP, RTSP, etc.
- ❑ Programming languages and frameworks for implementing dynamic content & behavior on client and server (e.g. JavaScript, jQuery, Backbone, XSLT, Java, JSP, PHP, RoR, SQL)
- ❑ W3C (World-Wide-Web Consortium, <http://www.w3.org>) established to promulgate Web technology standards (do standards really matter?)

Web-Internet Relationship

- ❑ From the Internet's point of view, the Web's HTTP protocol is just one among many application protocols, along with SMTP (email), FTP, SIP, DNS, etc.
- ❑ Internet is designed as a layered architecture whose layers include:
 - LAN (Local Area Network) connects hosts, e.g. with Ethernet (link-layer address, e.g.: 00-B0-D0-3E-51-BC)
 - IP (Internet Protocol) for interconnecting separate physical networks (LANs) uses IP addresses, e.g.: 142.150.160.47)
 - TCP (Transmission Control Protocol) for reliably sending streams of bytes to TCP ports on a host on an IP network
 - Applications, such as the Web, running on network-attached hosts, exchange information over the above layers
 - DNS (Domain Name Service) is a special application protocol for translating names and domains into IP addresses (e.g. `www.utsc.utoronto.ca == 142.1.96.30`)

Internet Protocol stack



Internet's Role in Web Operation

- ❑ The Internet is a fascinating system in its own right, but we have time only for a snapshot overview:
 - The Internet Protocol (IP) provides unreliable best-effort packet-delivery service between host interfaces
 - The Transmission Control Protocol (TCP) provides a reliable ordered-byte-stream service on top of IP
 - Hypertext Transfer Protocol (HTTP) uses TCP to exchange messages between Web clients & servers
 - Domain Name System (DNS) uses the unreliable UDP protocol to provide an IP-address lookup service
- ❑ The gory details are covered in CSCD58

Internet Engineering & Governance

- ❑ Just as the Web has the W3C to manage its standardization process, the Internet has two standardization/governance bodies:
- ❑ Internet Engineering Task Force (IETF): as the name suggests this group is responsible for engineering the Internet protocols, including HTTP (Web), TCP (generic Internet reliable connection service), IP (the unreliable Internet packet protocol) and many others
- ❑ Internet Corporation for Assigned Names and Numbers (ICANN) has a more politicized role, in managing the distribution of IP address blocks to regions, countries and ISP's, as well as the process of managing the Internet's domain name system



Client-Server Architecture

- ❑ Web Browsers (clients)
 - Chrome (the C09 reference browser), Firefox, IE, Safari, Opera, Blackberry, Lynx, ...
- ❑ Web Servers
 - Apache, njinx, Tomcat, Microsoft IIS, Node.js ...
- ❑ Naming via URLs (Uniform Resource Locators)
 - `<scheme>:<scheme-specific-part>`, e.g.
 - `http://<fully-qualified-domain-name>/<path>`, e.g.
 - `http://www.utsc.utoronto.ca/~rosselet/csc09/`
- ❑ HTTP (HyperText Transfer Protocol) and MIME (Multipurpose Internet Mail Extensions – why need MIME?)
 - `<type>/<subtype>` e.g.: `text/html`, `text/xml`, `image/png`, `audio/mp4`, `video/H261`, `application/json` ...

Static-Content Case

- ❑ Type a URL into a browser window, e.g.:
 - `http://www.utoronto.ca`
- ❑ Since the Internet's plumbing works with numbers not names, DNS is invoked to convert to numeric host IP number: `142.1.96.30`
- ❑ Browser makes a TCP connection to "port" 80 on the server and sends an HTTP GET request (why 80?)
- ❑ Web Server running on `www.utoronto.ca` fetches the default page for this site and sends it back to the browser via TCP and the HTTP protocol
 - response header provides info to help the browser interpret the page coming from the server: the content type (e.g. MIME text/html), character encoding (e.g. ISO-8859-1), etc.
- ❑ Server may elect to forward or redirect the request

```
> telnet www.utsc.utoronto.ca 80
Trying 142.1.96.30...
Connected to www.utsc.utoronto.ca.
Escape character is '^['.
```

Try it!

(can be useful as a debugging aid, to see unprocessed details of a server's response)

```
GET /. HTTP/1.0 \n\n
```

```
HTTP/1.1 200 OK
Server: nginx
Date: Thu, 28 Aug 2014 18:39:38 GMT
Content-Type: text/html
Connection: close
Set-Cookie: Apache=142.1.96.164.1409251178636876; path=/
Accept-Ranges: bytes
Cache-Control: no-cache,must-revalidate
Pragma: no-cache
<!doctype html>
```

```
<html xmlns="http://www.w3.org/1999/xhtml" lang="en" xmlns:html5="http://www.w3.org/1999/xhtml" >
<head>
<meta http-equiv="X-UA-Compatible" content="IE=EmulateIE7, IE=edge, chrome=1" />
<meta name="language" content="en" />

<meta charset="utf-8" />

<meta name="viewport" content="width=device-width; initial-scale=1"/> <!-- very important 4
mobile devices -->

<meta name="University of Toronto Scarborough" content="The homepage of the University of Toronto
Scarborough"/>
<meta name="description" content="The home page of the University of Toronto Scarborough."/>
<meta name="keywords" content="University of Toronto at Scarborough"/>
...
```

Dynamic-Client Case

- ❑ Interaction between client and server similar to static case, but response page contains code that executes within the browser, e.g.
 - JavaScript
 - Flash, Silverlight etc. plugins
- ❑ Ajax works with JavaScript to make interactive server requests (that do not update entire page)
 - Mechanism behind much of “Web 2.0”, “RIA”, “SPA”
- ❑ Security always an issue with dynamic client content
 - Need to keep client-host safe from malicious code running in the browser (that could change/copy/send information on the client system, install malware, etc.)
 - Web-based malware, has become a major irritant/threat

Dynamically-Served Content Case

- ❑ Page requests trigger code-execution on server side
 - Usually the output produced by this code is returned to the browser – e.g. code could trigger a search of a server-side persistent data store whose results would be returned in the response
- ❑ Widely used technologies:
 - Common Gateway Interface (CGI) can be used with a wide variety of server languages)
 - Server-Side Includes (SSI)
 - JSP, ASP, PHP, RoR, etc. embed server commands w/i HTML docs
 - Java Servlets/JSP execute on Web server
 - RESTful API's – language agnostic, typically used with “thinner” servers
- ❑ Security a concern here too, but now the issue shifts to server-side security (e.g. poorly written CGI could open door to malicious actions against server)

Browser-Document Compatibility

- ❑ Backward compatible (browser looking backward)
 - A new browser renders an earlier-generation document
 - ❑ deprecated features and older syntax must live on in browsers
- ❑ Forward compatible (browser looking forward)
 - An older browser renders a later-generation document
 - ❑ requires hacks in the document; author includes support for older browsers (or in Google's case, renders old UI for user)
- ❑ Sideways compatible
 - Browsers from different companies render documents in a consistent way
 - ❑ requires standards (the W3C's role)
 - ❑ although standards are widely implemented, still often must use document hacks (e.g. for CSS, JavaScript, DOM, Ajax)

Web Synopsis

- ❑ Invented by Tim Berners-Lee, drawing upon a number of existing technologies/systems including: SGML, hypertext, the Internet, Unix, client-server architecture
- ❑ TBL designed and implemented HTML, networked hyperlinks, URL's, HTTP
- ❑ The Web consists of a collection of markup and style languages, network protocols, programming languages
- ❑ Web standards are set by the W3C (and why stds matter)
- ❑ Internet serves as Web communication system
- ❑ Browsers and servers operate in a client-server relationship:
 - HTTP and MIME protocols
 - Static/dynamic client, server scenarios
- ❑ Browser compatibility is an issue for Web apps