

CSCC09F

Programming on the Web



Protocols and HTTP, Web Browsers and Servers, Web Caching

first some background on network-based
applications ...

Client-Server Paradigm


Client:

- ❑ initiates contact with server (“speaks first”)
- ❑ typically requests service from server
- ❑ for Web, client is implemented in browser; for e-mail, in mail user agent

Server:

- ❑ provides requested service to client
- ❑ e.g., Web server sends requested Web page, mail server delivers e-mail

Typical network app has two pieces: client and server




Applications and application-layer protocols

Application: communicating, distributed processes

- running in network hosts in “user space”
- exchange messages to implement application
- e.g., email, file transfer, the Web, network news

Application-layer protocols

- one “piece” of an app
- define messages exchanged by apps and actions taken
- uses services provided by lower layer protocols



HyperText Transfer Protocol (HTTP)

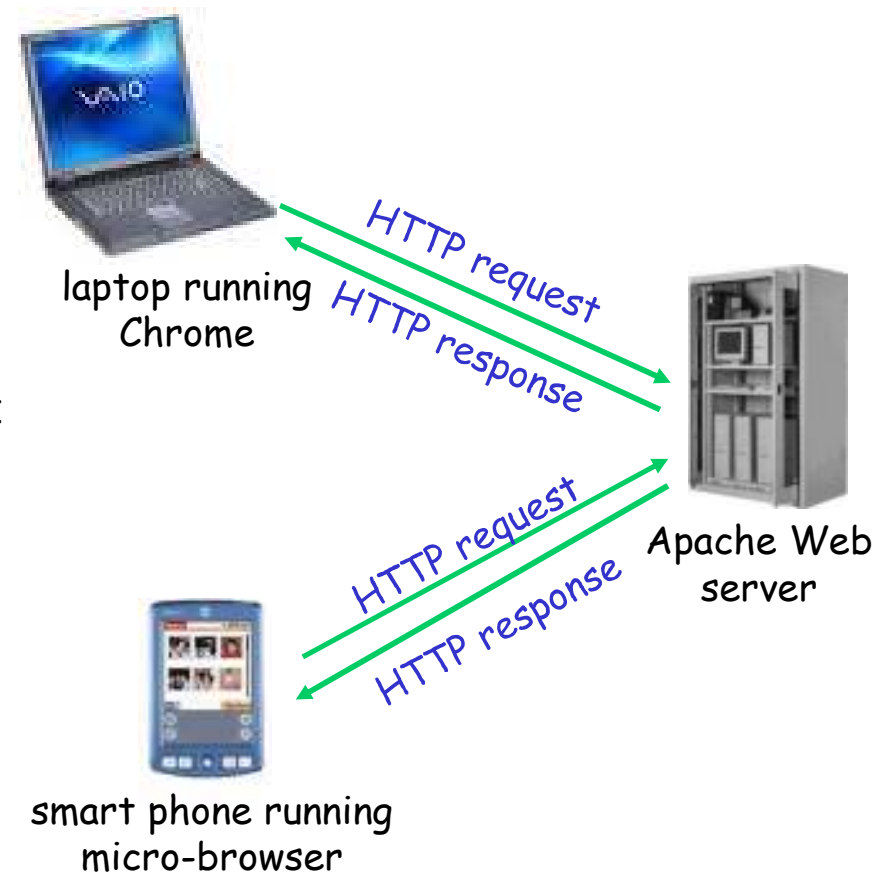
Part of Tim Berners-Lee's innovation

- Share documents across the network
- Defined 1989-1991
- Standardized and much expanded by the IETF
- Rides on top of “TCP” transport protocol
 - TCP provides: reliable, bi-directional, in-order byte stream (exists to compensate for unreliable Internet Protocol IP)
- Goal: transfer objects between systems
 - Do not confuse with other WWW concepts:
 - HTTP is not a page-layout language (that is HTML)
 - HTTP is not an object-naming scheme (that is URLs)

The Web: HTTP protocol

HTTP: HyperText Transfer Protocol

- ❑ Web's application-layer protocol
- ❑ client/server model
 - **client**: (active) browser that requests, receives, "displays" Web objects
 - **server**: (passive) Web server sends objects in response to requests
- ❑ http1.0: RFC 1945
- ❑ http1.1: RFC 2068



HTTP 1.0

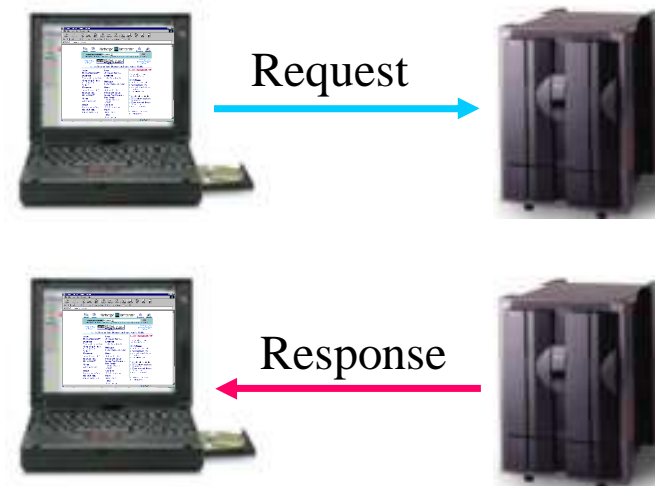
- ❑ Interaction between Web client (browser) and Web server occurs in two phases:

- **Request phase:**

- Browser requests page from Web server

- **Response phase:**

- Server sends back requested page or code



- ❑ Protocol for each phase consists of:

- **Header** (2 parts: request line or response status (single line), header fields (possibly multiple lines))
 - **Body** (request or response)

Try out HTTP (client side) for yourself

1. Telnet to your favorite Web server, e.g.:

```
telnet www.utsc.utoronto.ca 80
```

Open TCP connection to port 80 (default http server port) at www.utsc.utoronto.ca. Anything typed is sent to port 80 at www.utsc.utoronto.ca

2. Type in a GET HTTP request:

```
GET /~rosselet/csc09/ HTTP/1.0
```

Type this at the prompt (followed by two carriage returns) to send this minimal GET request to the HTTP server

3. Observe response message sent by http server

```
mathlab:~> telnet www.tdcanadatrust.com 80
```

```
Trying 184.86.36.40...
```

```
Connected to e9726.b.akamaiedge.net.
```

```
Escape character is '^['.
```

```
GET /index.html HTTP/1.0
```

```
Host: www.tdcanadatrust.com
```

```
HTTP/1.0 301 Moved Permanently
```

```
Location: http://www.tdcanadatrust.com/products-  
services/banking/index-banking.jsp
```

```
Content-Length: 280
```

```
Content-Type: text/html; charset=iso-8859-1
```

```
Date: Wed, 22 Oct 2014 12:03:45 GMT
```

```
Connection: close
```

```
<!DOCTYPE HTML PUBLIC "-//IETF//DTD HTML 2.0//EN">
```

```
<html><head><title>301 Moved Permanently</title></head><body>
```

```
<h1>Moved Permanently</h1>
```

```
<p>The document has moved <a
```

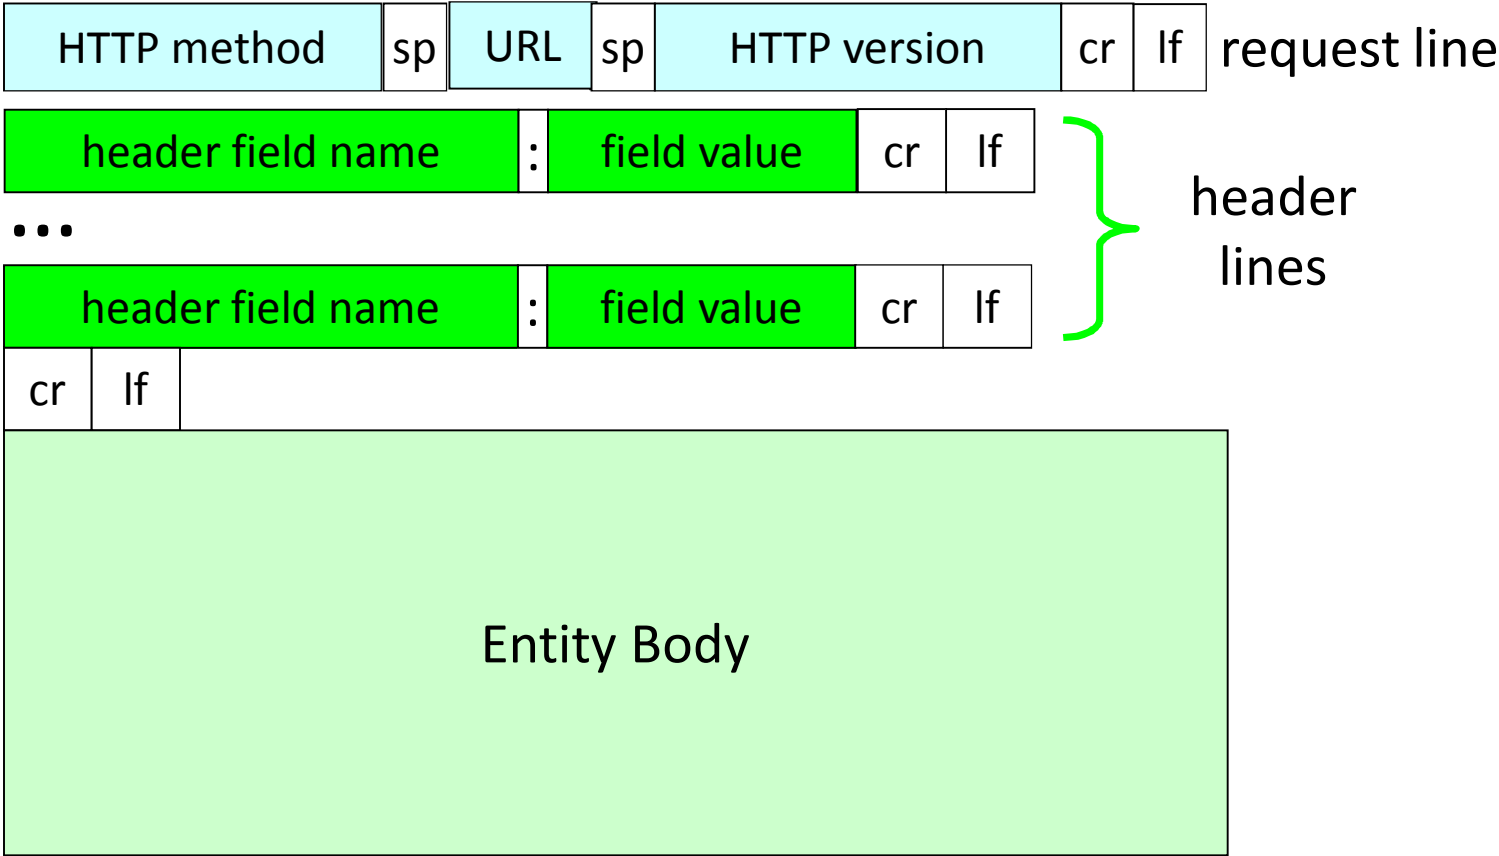
```
href="http://www.tdcanadatrust.com/products-  
services/banking/index-banking.jsp">here</a>.</p>
```

```
</body></html>
```

```
Connection closed by foreign host.
```

Example: command line HTTP request

HTTP request message format



HTTP 1.0 Request Phase

□ Request line

HTTP method	sp	URL	sp	HTTP version	cr	lf
-------------	----	-----	----	--------------	----	----

□ HTTP method

- GET – return content of specified document
- HEAD – return only headers of GET response
- POST – execute specified doc with enclosed data

□ URL

- Usually expressed as: /path (different for HTTP/1.1)
 - e.g. /index.html

□ HTTP version

- e.g. HTTP/1.0

HTTP 1.0 Request Phase (cont)

- Header fields, sequence of:

header field name	:	field value
-------------------	---	-------------

- Examples:

- `Accept: text/html`
- `Accept: image/jpg`
- `Accept-language: en; fr`
- `If-modified-since: 20 Oct 2014`
- `Content-Length: 7814`

HTTP request message example

□ http request message:

- Note all HTTP messages are in ASCII (text format)

```
GET /appdir/page.html HTTP/1.1
User-agent: Chrome/37.0.2062.124
Accept: text/html,image/gif,image/jpeg
Accept-language:fr
```

```
data data data data data ...
data data data data data ...
data data data data data ...
```

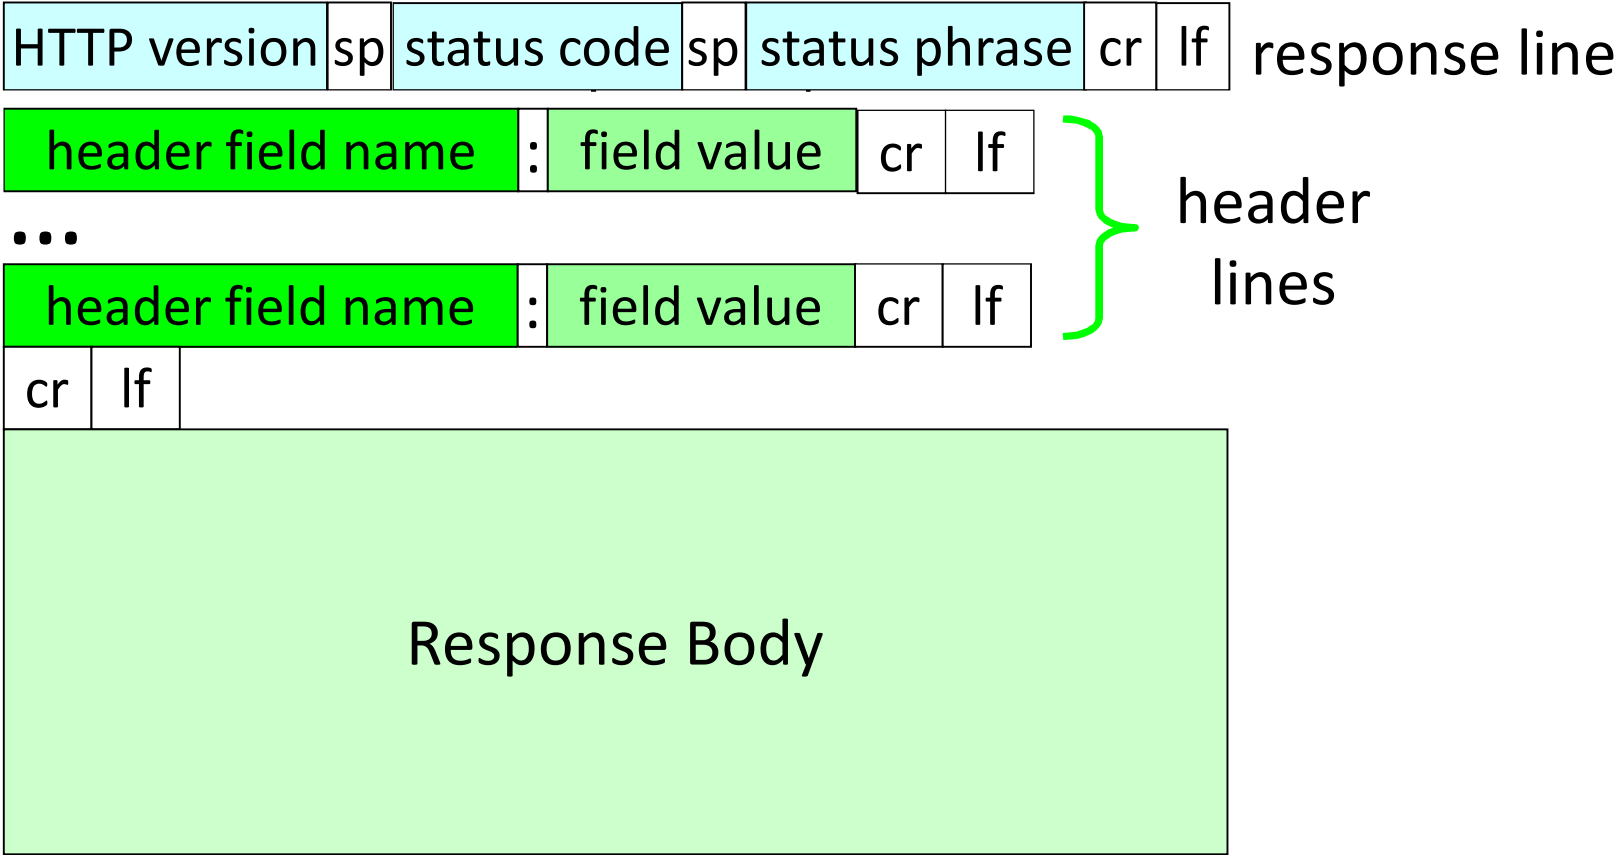
request line
(GET, POST,
HEAD, etc.,
commands)

header
lines

blank line
indicates end
of header

Post method data,
e.g.,
form field values

HTTP response message format



HTTP 1.0 Response Phase

- ❑ Response consists of
 - ❑ Status line:
 - ❑ HTTP version
 - ❑ 3-digit status code (success, error, redirection, etc)
 - ❑ Brief text explanation of status code (e.g. OK)
 - ❑ Response Header fields:
 - ❑ Other page attributes (content type, content length, expiration, last modified, server type, etc)
 - ❑ Additional information (if redirection, other location)
 - ❑ empty line (delimiter between header and body)
 - ❑ Response body (e.g. html for requested page)

Response Status Codes

■ 3 digit response code

- 1XX – informational
- 2XX – success (return requested document)
- 3XX – redirection (send request elsewhere)
- 4XX – client error (e.g. document not found)
- 5XX – server error (e.g. Tomcat setup problem)

■ status phrase – brief translation of the numeric response code (e.g. 404 – File not found)



Response Status Code Examples

□ Some commonly occurring status codes:

200 OK

- request succeeded, requested object appears later in this message

301 Moved Permanently

- requested object moved, new location specified later in this message (Location:)

400 Bad Request

- request message not understood by server (malformed HTTP request)

404 Not Found

- requested document not found on this server

501 Internal Server Error

- something wrong on server side – e.g. CGI script error

HTTP response header

- ❑ A server may send a set of responses to a given client within a browsing session, e.g.: plain text, html, images, pdf files, etc.
 - ❑ How does the browser know what to do when a message from the server shows up?
 - ❑ content-type header is the key to identifying the message data-type and thus how the browser should handle it.
- »
- ❑ based on the MIME type-codes originally defined to allow attachments to be bundled with e-mail (RFC 1341)

HTTP response message example

HTTP/1.1 200 OK

Date: Weds, 22 Oct 2014 12:19:15 GMT

Server: Apache/2.2.22 (Ubuntu)

Last-Modified: Tue, 02 Sep 2014

Content-Length: 8740

Content-Type: text/html

data data data data data ...

data data data data data ...

data data data data data ...

status line:
(protocol,
status-code,
status-phrase)

header
lines

blank line
indicates end
of message
header

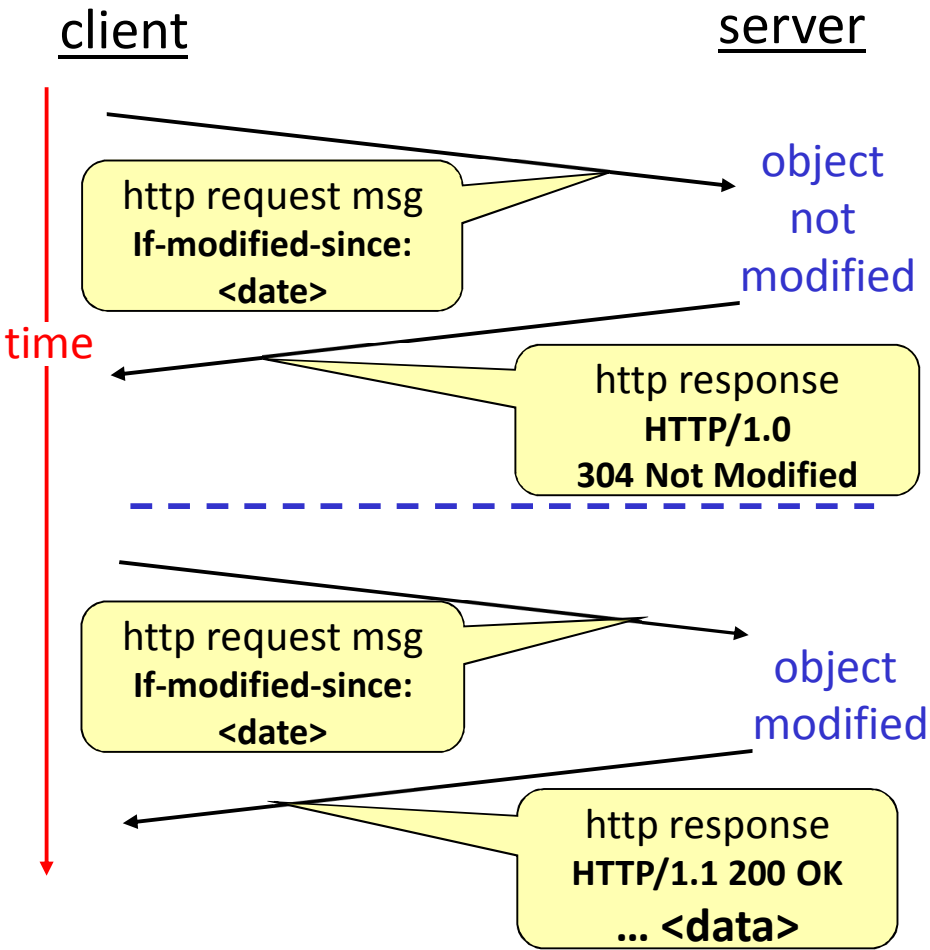
data, e.g.,
requested
html file

HTTP other features

- ❑ POST method
 - ❑ Client can send information to server, e.g. form field data, file-upload content
- ❑ Conditional GET **if-modified-since** request header
 - ❑ Client tells server it already has a copy of the requested data, and asks server whether this copy is stale version and thus fresh copy needs to be re-fetched from server
 - ❑ Why bother?
 - ❑ Efficiency: avoid (re)sending data if client already has up-to-date copy (cached)

User-server interaction: conditional GET

- ❑ client: specifies date of cached copy in request
If-modified-since:
<date>
- ❑ server: response contains no body if cached copy is up-to-date:
HTTP/1.0 304 Not Modified



HTTP Authentication

- ❑ Goal: controlled access to documents
- ❑ Basic authentication
 - When challenged, client sends user-id and password in clear to server
 - Not secure enough (snooping is easy) but useful for restricting access to low-value targets
- ❑ Protocol is stateless: client must present credentials in each request
 - but, browser caches user-id + password so user does not have to enter them for each request

Example:

```
private.html  
(user: csc09f14,  
pass: topsecret)
```

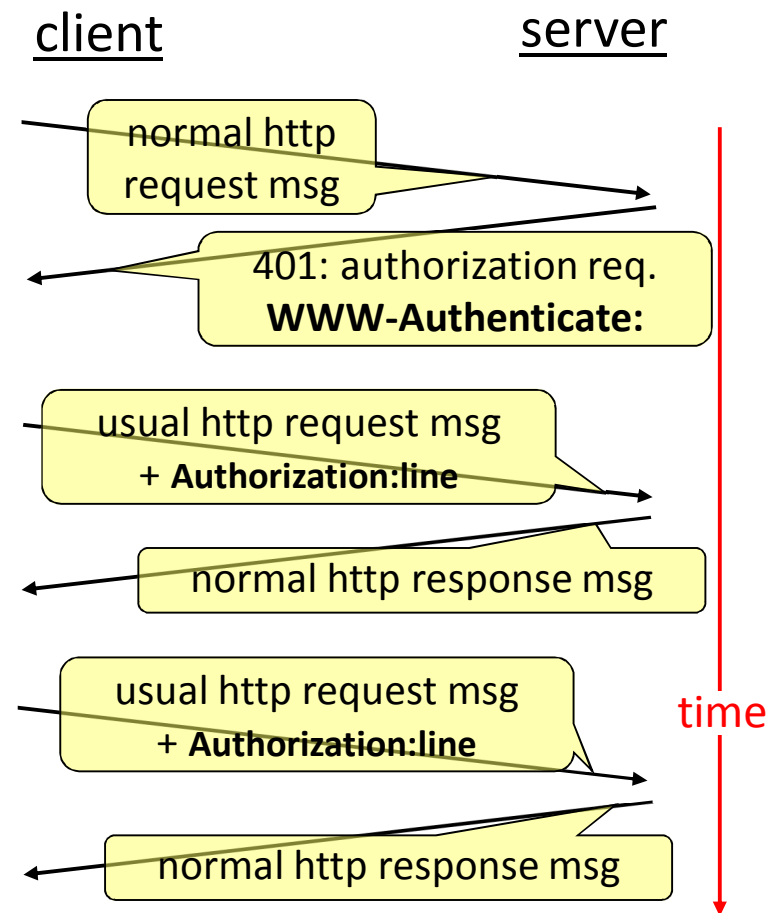
client-server interaction: authentication

- **WWW-Authenticate** response header: typically username, password based

- **Authorization:** request header line
- if no authorization presented, server refuses access, sends another

"WWW authenticate:"
header line in response

- Set up with **.htaccess**, **.htpasswd**



HTTP is “Stateless”

- ❑ Server maintains no information about past client requests – each one must be self-contained, why?
- ❑ Protocols that maintain “state” are complex – stateless simplifies protocol and server, e.g. server has nothing to recover after crash, smaller and simpler implementation
- ❑ Are there any implications for developers?
 - History (state) must be maintained at the application level – by apps themselves

Stateless? I don't get it.

- ❑ OK, so HTTP is stateless, but I don't quite see why that matters to the Web app developer
 - ❑ Think about how users interact with a Web application...
 - ❑ Typically they perform a sequence of steps, e.g. selecting items, making choices, and the user expects at the end of such a sequence that the application will know what they mean when they say “check out”.
 - ❑ But what is the server view of these user “steps”?
 - ❑ typically these are individual HTTP requests
 - ❑ But if the server does not maintain status information across client requests, then how will the server know what to do for “check out”?
 - ❑ Short answer, it won't. There's no way for the server to correlate multiple request from same user!

Cookies, who needs 'em?

- ❑ So, what's the answer? How can the server keep track of something when it has no concept of request "state"
- ❑ Solution: store a small amount of information about the client state in a cookie
 - ❑ Where to store cookies?
 - ❑ on client side! so how does the Server get them?
 - ❑ client sends cookie(s) with each request
 - ❑ what happens if the client is talking to multiple servers or if there are multiple cookies?
 - ❑ cookie values may be updated by server with response(s)

User-server interaction: cookies

- ❑ server sends “cookie” to client in response msg

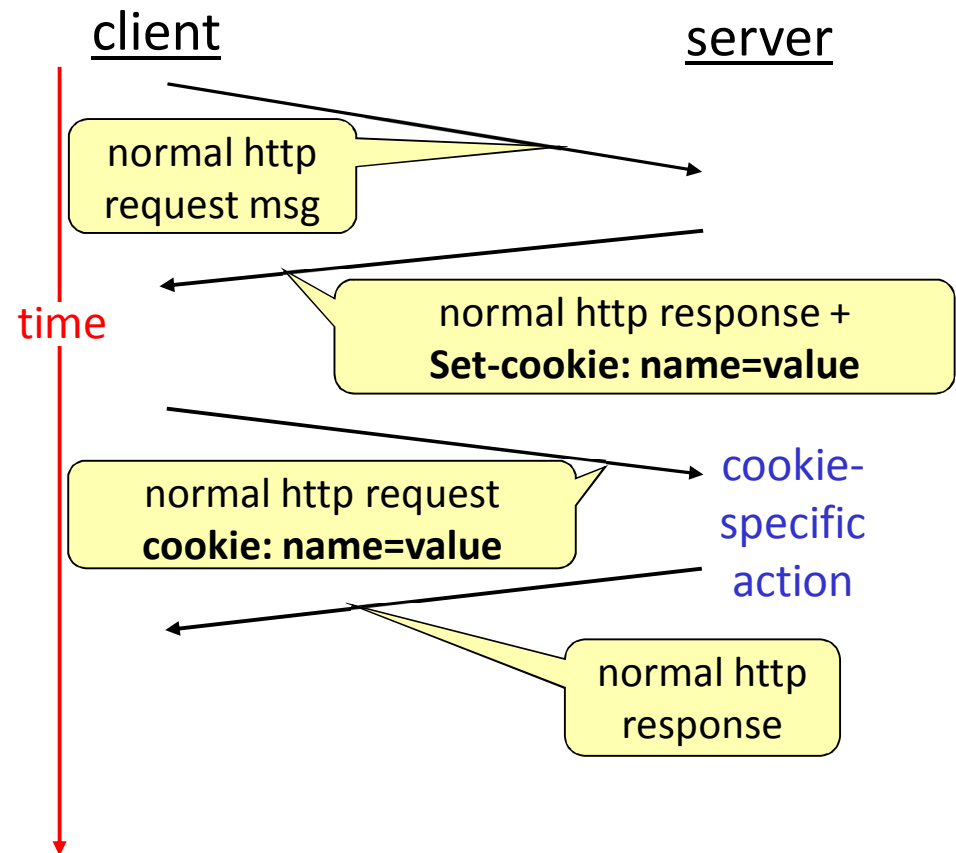
Set-cookie:
name=1678453

- ❑ client presents cookie in later requests

cookie: name=1678453

- ❑ server may match presented-cookie with server-stored info.

- ❑ uses:
 - authentication
 - remembering user preferences, previous user selections



Cookies, what's in em?

- ❑ How much client state can you store in a cookie?
- ❑ Is it a good idea to have client information going back and forth between client and server?
 - What if something gets lost, or changed, does that mean the state changes?
 - What if somebody else can read the cookies from the network, how is user privacy protected?
 - Although cookies can store a modest amount of data (on the order of 4KB, and with up to 20 cookies per server), generally that's not how they are used.
 - Instead, the actual data is stored on the server, e.g. in a DB, and the cookie holds just a random-looking “token” value that allows the server to locate the relevant data.
 - This approach solves both of the above concerns


Web Caching

- ❑ Goal: fill client (browser) request without going to originating server. Why bother?
- ❑ In early days of Web, the Internet was a lot slower – mask this by retrieving content from local copies.
- ❑ Internet faster now, but still pays to be efficient in use of “bandwidth” – network capacity.
- ❑ We’ve already seen one HTTP request header designed for this purpose; which one?
- ❑ Client implements its own caching mechanism.
- ❑ Important to recognize that a “Web page” may consist of multiple objects, each of which is transferred/cached separately, esp. images.

Proxy Server Web Caches


Goal: fill client request more efficiently than sending request all the way to originating server

- ❑ user configures browser to use proxy-server
- ❑ client sends all HTTP requests to proxy server
 - if object exists in proxy's cache, proxy immediately returns object in http response
 - else proxy requests object from origin server, then returns http response to client and retains cached copy



Benefits of Web Caching

- Assume:** cache is “close” to client (e.g., in same local area network)
- ❑ smaller response time: cache “closer” to client than origin is
 - ❑ decrease traffic to distant servers
 - link out of institutional/local ISP network often bottleneck



Caching, why?

- ❑ improves performance for users:
 - Response time
 - Availability
 - Scalability
 - Load balancing
- ❑ conserves network and server resources
- ❑ requirement: logical transparency:
 - the presence of a proxy-cache between client and server should not change the information content received by the client in response to requests, e.g. client should never get stale data